

## A

- Acceptable Use Policy 479
- access control lists (ACLs) 189, 193
- access control, database
  - network access control 310–311
  - overview 311
  - using IAM service 309–310
- account, AWS
  - creating
    - choosing support plan 26
    - contact information 24
    - creating key pair 29–32
    - login credentials 23
    - payment details 24
    - signing in 28–29
    - verifying identity 25
  - security
    - authentication roles 177–179
    - authorization policies 174–176
    - creating users 176–177
    - IAM service 173–174
    - root user 172–173
- ACID (atomicity, consistency, isolation, and durability) 294, 349
- ACLs (access control lists) 189, 193
- agent forwarding 187–188
- AKI (Amazon Kernel Image) 63
- alarms, CloudWatch 203, 387–388
- Amazon API Gateway 229
- Amazon EFS. *See* EFS (Elastic File System)
- Amazon ElastiCache. *See* ElastiCache
- AMI (Amazon Machine Image) 62–63, 139–140
- Apache Bench 479
- APIs (application programming interfaces) 102
- apt package manager 75
- archiving objects 242
- ARN (Amazon Resource Name) 175
- asynchronous decoupling
  - consuming messages 425–428
  - converting synchronous process to asynchronous 421–422
  - creating SQS queue 423
  - overview 420–421
  - sending messages to queue 423–425
  - SQS messaging limitations 428–429
  - URL2PNG application example 422–423
- atomic operations 255
- atomicity, consistency, isolation, and durability (ACID) 294, 349
- auto-scaling 397, 465
  - groups 397–398, 466
  - overview 396–407
  - triggering
    - based on CloudWatch metrics 472–474
    - based on schedule 471–472
    - overview 469–470
- automating operational tasks, with AWS Lambda
  - adding tag containing owner of EC2 instance
    - automatically 218–226
    - authorizing Lambda function to use other AWS services with IAM role 224–226
  - deploying Lambda function with SAM 226
  - implementing Lambda function in
    - Python 222–223
  - setting up Lambda function with SAM 223–224
  - subscribing to CloudWatch events 219–222
- building website health check 203–218
  - accessing endpoints within VPC 217–218
  - creating Lambda function 204–210

- automating operational tasks, with AWS Lambda
  - (*continued*)
  - monitoring Lambda function with CloudWatch metrics and alarms 212–216
  - using CloudWatch to search through Lambda function's logs 210–212
- executing code 200–203
  - comparing AWS Lambda with virtual machines (Amazon EC2) 202–203
  - running code on AWS Lambda 201–202
- Free Tier 228
- impacts of serverless pricing model 228–229
- limitations of 227–228
- pricing model 228
- use cases
  - data processing 230
  - IoT back end 231
  - web application 229
- AWS (Amazon Web Services)
  - advantages
    - automation capabilities 10
    - fast-growing platform 10
    - platform of services 10
    - reducing time to market 11
    - reliability 11
    - scalability 11
    - standards compliance 12
    - worldwide deployments 12
  - alternatives 15–16
  - costs 275
    - billing example 13–15
    - Free Tier 13
    - overview 12–13
    - pay-per-use pricing model 15
  - services overview 16–18
  - tools
    - blueprints 22
    - CLI 20
    - Management Console 19
    - SDKs 21
  - use cases
    - fault-tolerant systems 8–10
    - running Java EE applications 7–8
    - running web shop 5–7
- AWS Elastic Beanstalk 145
  - applications in Management Console 149–150
  - creating application 147
  - creating environment 147
  - deleting applications 150–151
  - deployment 457–461
  - deployment options comparison 138–139
  - describing status of installation 147–148
  - overview 145–146
  - uploading Zip archive 147
- AWS Lambda
  - adding tag containing owner of EC2 instance automatically 218–226
  - authorizing Lambda function to use other AWS services with IAM role 224–226
  - deploying Lambda function with SAM 226
  - implementing Lambda function in Python 222–223
  - setting up Lambda function with SAM 223–224
  - subscribing to CloudWatch events 219–222
- building website health check with 203–218
  - accessing endpoints within VPC 217–218
  - creating Lambda function 204–210
  - monitoring Lambda function with CloudWatch metrics and alarms 212–216
  - use CloudWatch to search through Lambda function's logs 210–212
- executing code with 200–203
  - comparing AWS Lambda with virtual machines (Amazon EC2) 202–203
  - running code on AWS Lambda 201–202
- Free Tier 228
- impacts of serverless pricing model 228–229
- limitations 227–228
- pricing model 228
- use cases
  - data processing 230
  - IoT back end 231
  - web application 229
- AWS OpsWorks
  - deployment options comparison 138–139
  - multilayer applications using
    - accessing kiwiIRC 162–163
    - adding app to Node.js layer 160–161
    - adding client and server instances for IRC chat 161
    - creating custom layers 156–159
    - creating Node.js layer 155–156
    - creating stack 154
    - deleting stacks 163–164
  - overview 151–153
- AZ (availability zones) 274, 314, 407
  - defined 392
  - outages 392–396
    - auto-scaling 396–399
    - IP addresses and 407–411
    - network-attached storage and 402
    - recovering failed EC2 instance to another availability zone 399–407
  - recovering failed EC2 instance to another availability zone 399–407
  - redundancy 436

**B**

backup 289–292  
     automated snapshots 303–304  
     copying database to other region 307  
     manual snapshots 304–305  
     restoring database 305–307  
 !Base64 function 142  
 bastion host  
     defined 186  
 benchmarking 479  
 block-level storage  
     instance stores  
         backups 272  
         overview 268–271  
         performance testing 272  
         viewing and mounting volumes 271–272  
     network-attached storage  
         backing up data 266–268  
         creating volumes 260–261  
         managing volumes 261–263  
         performance improvements 263  
 BlockDeviceMappings 269–270  
 blueprints  
     AWSTemplateFormatVersion value 123  
     example template 127–133  
     Lambda function configuration 205  
     outputs structure 126  
     overview 22, 121–122  
     parameters structure 123–125  
     resources structure 125–126  
 bucket policies 253  
 BucketNotEmpty error 240  
 buckets, S3  
     configuring 253–254  
     creating programmatically 253  
     linking custom domain 255  
     listing files 251–252  
     setting up 249  
     versioning 239–240  
 burst rate, EFS 286  
 burstable performance 474  
 BurstCreditBalance metric 286–287

**C**

Cache Engine 334  
 cache, invalidating 439  
 caching data in memory. *See* ElastiCache  
 calculator for monthly costs 13  
 Cassandra 354  
 CDN (content delivery network) 6, 83, 252  
 Chef 138, 152  
 CIDR (Classless Inter-Domain Routing) 185

CLI (command-line interface) 237  
     –query option 114–115  
     advantages of using scripts 108–109  
     configuring user authentication 110–113  
     creating virtual machine using script 113–115  
     help keyword 113  
     installing  
         on Linux 109  
         on Mac OS 109  
         on Windows 109–110  
     listing EC2 instances 113  
     overview 20  
     usage overview 113  
 cloud computing  
     deployment environment 136–137  
     overview 4–5  
 CloudFormation  
     alternatives 133  
     blueprints  
         AWSTemplateFormatVersion value 123  
         example template 127–133  
         outputs structure 126  
         overview 121–122  
         parameters structure 123–125  
         resources structure 125–126  
 CloudFormation template  
     defining Lambda function with SAM within 224  
     minimal 327–328  
 CloudWatch  
     creating alarm 387–388  
     monitoring Lambda function with metrics and  
         alarms 212–216  
     overview 385–386  
     subscribing to events 219–222  
     triggering auto-scaling based on metric 472–474  
 CNAME records 255  
 cold-start 227  
 content delivery network (CDN) 6, 83, 252–253  
 cookbooks, Chef 152  
 cost  
     AWS Lambda 203  
     billing example 13–15  
     DynamoDB service 352  
     Free Tier 13  
     optimizing for virtual machines  
         overview 92  
         reserved virtual machines 93–94  
         spot instances 95–100  
     overview 12–13  
     pay-per-use pricing model 15  
 CPU load 474  
 CPUUtilization metric 344–345  
 CREATE\_COMPLETE state 43, 298–299, 479  
 CREATE\_IN\_PROGRESS state 43  
 CRUD (create, remove, update, delete) 248

**D**


---

data centers  
   locations 4, 12  
   starting virtual machine in different 82–86

data processing, automating operational tasks with  
   AWS Lambda 230

databases  
   for WordPress blogs example 49–51  
   network access control 310–311

dd utility 263

DDoS (Distributed Denial of Service) attacks 167

declarative approach, defined 121

decoupling  
   asynchronous, with message queues  
     consuming messages 425–428  
     converting synchronous process to  
       asynchronous 421–422  
     creating SQS queue 423  
     overview 420–421  
     sending messages to queue 423–425  
     SQS messaging limitations 428–429  
     URL2PNG application example 422–423

concept explained 413–415

dynamic EC2 instance pools  
   by load balancer 476–480  
   by queue 481–484  
   overview 475–482

redundant EC2 instances 436–437

synchronous, with load balancers  
   overview 415–416  
   setting up load balancer 416–420

default region setting 39

dependencies, CloudFormation templates 122

deployment  
   AWS Lambda 203  
   comparison of options 138–139  
   defined 135  
   in scalable cloud environment 136–137  
   multilayer applications with AWS OpsWorks  
     Stacks 152–153  
   running script on virtual machine startup  
     overview 139  
     using user data 140  
   web applications with AWS Elastic Beanstalk  
     components 146  
     creating application 147  
     creating environment 147  
     deleting applications 150–151  
     uploading Zip archive 147

DevOps (development operations)  
   movement 104–105

df command 262

disaster recovery 411–412

Discourse application, installing with  
   CloudFormation 336–344  
   cache 338–339  
   database 339–340  
   testing CloudFormation template 342–344  
   virtual machine 340–342  
   VPC 337–338

Distributed Denial of Service (DDoS) attacks 167

DLQ (dead-letter queue) 456

DNS (Domain Name System) 83

DSL (domain-specific language) 152

DSS (data security standard) 12

dynamic EC2 instance pools  
   decoupling  
     by load balancer 476–480  
     by queue 481–484  
     overview 475–482  
   managing 465–469

DynamoDB service  
   administrative tasks 352  
   costs 352  
   deleting data 373–374  
   eventual consistency 372–373  
   modifying data 374–375  
   NoSQL comparison 354  
   overview 349–351  
   primary key  
     defined 355  
     partition key and sort key 355–356  
   querying data  
     by key 365–366  
     by key and filter 366–368  
     overview 364–365  
     retrieving all items 371–372  
     using secondary indexes 368–371

RDS vs. 353

running DynamoDB locally 356

scaling 375–379

tables  
   creating 358  
   overview 354–355  
   using partition key 358–359  
   using partition key and sort key 360–361

**E**


---

EBS (Elastic Block Store)  
   backing up data from 266–268  
   creating volumes 260–261  
   defined 432  
   managing volumes 261–263  
   performance improvements 263

EBS. *See* instance stores

EC2 (Elastic Compute Cloud) service  
   comparing AWS Lambda with 202–203  
   defined 3, 37  
   recovering instances 386  
   *See also* virtual machines

EFS (Elastic File System)  
   backing up data 289–292  
     using CloudFormation to describe EBS volume 290  
     using EBS volume 290–292  
   burst rate 286  
   filesystem, creating 277–278  
     pricing 277–278  
     using CloudFormation to describe filesystem 277  
   filesystem, monitoring 286–289  
     Max I/O Performance Mode, whether to use 286  
     monitoring permitted throughput 287–288  
     monitoring usage 288–289  
   mount target, creating 278–279  
   mounting EFS share on EC2 instances 280–283  
   performance, tweaking 284–286  
     expected throughput 285–286  
     Performance mode 285  
   sharing files between EC2 instances 283–284

Elastic Beanstalk. *See* AWS Elastic Beanstalk

elastic IP addresses 86

ElastiCache  
   access to cache, controlling 334–336  
     access to configuration 334  
     cluster and data access 335–336  
     network access 334–335  
   cache deployment options 330–334  
     Memcached 330–331  
     Redis 331–334  
   creating cache cluster 327–330  
     minimal CloudFormation template 327–328  
     testing Redis cluster 328–330  
   monitoring cache 344–346  
     monitoring host-level metrics 344–345  
     Redis replication, whether up-to-date 345–346  
     sufficient memory 345  
   performance of cache, tweaking 346–348  
     compressing data 348  
     selecting right cache node type 347  
     selecting right deployment option 347–348

elasticity 464

ELB (Elastic Load Balancing) 37, 415

ENI (Elastic Network Interface) 432

environment variables  
   temporary credentials and 224  
   to pass settings to Lambda function 209

environment, defined 458

Erlang programming language 437

Errors metric, CloudWatch 212

eventual consistency 255, 372–373, 439

Evictions metric 344–345

## F

---

failure recovery  
   availability zone outages  
     auto-scaling 396–399  
     availability zones 392–396  
     IP addresses and 407–411  
     network-attached storage and 402  
     recovering failed EC2 instance to another availability zone 399–407

fault-tolerance  
   code considerations  
     “let it crash” concept 437–438  
     idempotent retry 438–439  
   defined 432  
   high availability vs. 384  
   overview 431–433  
   redundant EC2 instances  
     decoupling required 436–437  
     overview 434  
     removing single point of failure 434–436

web application  
   creating process 446–448  
   idempotent state machine 443–444  
   idempotent state transitions 444  
   Imagery application overview 440–461  
   looking up process 448–449  
   server 445–446  
   uploading files 449–452  
   worker for consuming SQS messages 452–455

web application deployment  
   DynamoDB 455–457  
   Elastic Beanstalk for worker 459–461  
   IAM roles 457  
   overview 455  
   S3 455–457  
   SQS 455–457

fdisk utility 261, 271

filesystem. *See* EFS (Elastic File System)  
   *See also* EBS; instance stores

firewalls 69, 158

Free Tier 13

fsfreeze command 267

## G

---

generations, defined 65

GiB (gibibyte) 265

Glacier service  
   adding lifecycle rule to bucket 242  
   moving objects to and from 245–248  
   S3 service 241  
 globally unique identifiers 236–237

**H**

---

HA (high availability)  
   defined 384, 432  
   disaster-recovery requirements 411–412  
   fault tolerance vs. 384  
   for databases 311–314  
   recovering from availability zone outages  
     auto-scaling 396–399  
     availability zones 392–396  
     IP addresses and 407–411  
     network-attached storage and 402  
     recovering failed EC2 instance to another  
       availability zone 399–407  
   recovering from EC2 instance failure with  
     CloudWatch  
     creating alarm 387–388  
     overview 385–386  
   redundant EC2 instances for  
     decoupling required for 436–437  
     overview 434  
     removing single point of failure 434–436  
 host machines 60  
 host-level metrics 344–345  
 httpd-tools package 479  
 HVM (Hardware Virtual Machine) 64

**I**

---

IaaS (infrastructure as a service) 5, 104  
 IAM (Identity and Access Management)  
   service 173–174, 309, 334  
   database access control using 309–310  
   roles  
     authorizing Lambda function to use AWS  
       services with 224–226  
     deployment 457  
 ICMP (Internet Control Message Protocol)  
   182–183  
 idempotent  
   defined 429, 438  
   retry 438–439  
   state machine 443–444  
   state transitions 444  
 IGW (internet gateway) 190–191  
 infrastructure as code  
   CLI  
     advantages of using scripts 108–109

  configuring user authentication 110–113  
   creating virtual machine using script 113–115  
   help keyword 113  
   installing 109–110  
   listing EC2 instances 113  
   usage overview 113  
   defined 104  
   DevOps movement 104–105  
   JIML 105–108  
   using SDKs 117  
 inline policy 176  
 installation  
   software on virtual machines 75  
 instance family groups 65  
 instance stores  
   backups and 272  
   overview 268–271  
   performance testing 272  
   viewing and mounting volumes 271–272  
 instances, defined 152  
 IOPS (input/output operations per second)  
   264–266  
 IoT 231  
 IP (Internet Protocol) 180  
 IP addresses  
   allocating for virtual machine 86–88  
   public vs. private 184  
 IRC (Internet Relay Chat) 153

**J**

---

Java EE applications 7–8  
 JIML (JSON Infrastructure Markup  
   Language) 105–108  
 JMESPath 114  
 jump boxes 186

**K**

---

key pair for SSH  
   creating 29–32  
   selecting for virtual machine 70–72  
 key-value stores 325, 350, 364  
 kiwiIRC 154–161, 163

**L**

---

Lambda function 203  
   configuring scheduled event triggering 208  
   creating 205–206  
   defining with SAM within CloudFormation  
     template 224  
   invocations 210  
   log messages from 211

- Lambda function (*continued*)
  - monitoring with CloudWatch metrics and alarms 212
  - website health check and environment variables to pass settings to 209
- Lambda. *See* AWS Lambda
- launch configurations 397
- lazy-loading 322
- “let it crash” concept 437–438
- lifecycle rules 242
- linkchecker tool 75
- Linux
  - connecting to virtual machines from 73
  - installing CLI on 109
- load balancers
  - decoupling dynamic EC2 instance pools by 476–480
  - for WordPress blogs example 47–48
  - synchronous decoupling with overview 415–416
  - setting up load balancer 416–420
- load monitoring 77–78
- load tests 479
- logs
  - viewing for AWS Elastic Beanstalk application 150
  - viewing for virtual machines 76–77

## M

---

- Mac OS X
  - connecting to virtual machines from 73
  - installing CLI on 109
  - key file permissions 31
- managed policy 176
- Management Console
  - overview 19
  - signing in 28
- Max I/O Performance Mode 286
- Memcached 326, 330–331
- memory, caching data in. *See* ElastiCache
- message queues
  - consuming messages 425–428
  - converting synchronous process to asynchronous 421–422
  - creating SQS queue 423
  - overview 420–421
  - sending messages to queue 423–425
  - SQS messaging limitations 428–429
  - URL2PNG application example 422–423
- metadata 236
- MFA (multifactor authentication) 172
- mkfs command 262
- MongoDB 354

- monitoring
  - load 77–78
  - viewing logs 76–77
- MySQL databases
  - costs 300
  - database instance information 299–300
  - exporting 301
  - for WordPress blogs example 49–51
  - WordPress platform using 297–299

## N

---

- Name tag 67
- NAT (Network Address Translation) 7, 184, 190, 196–197
- Neo4j 354
- network access control 310–311
- network-attached storage
  - availability zone outages and 402
  - backing up data from 266–268
  - creating volumes 260–261
  - managing volumes 261–263
  - performance improvements 263
- networking
  - ACLs (access control lists) 189, 193
  - controlling traffic 183–189
    - overview 179–181
    - using security groups 181–182
  - for virtual machines
    - allocating fixed public IP address 86–88
    - creating additional network interface 88–92
    - using access control lists (ACLs) 189, 193
- Node.js
  - installing 117, 249, 361, 423
- NoSQL databases 229
  - RDS vs. DynamoDB 353
  - running DynamoDB locally 356
  - scaling 375–379
  - tables
    - creating 358
    - overview 354–355
    - using partition key 358–359
    - using partition key and sort key 360–361

## O

---

- object stores 229
  - backing up data using 238–240
  - concepts of 236–237
  - data consistency 255
  - Glacier service
    - adding lifecycle rule to bucket 242
    - creating S3 bucket for 241–242

- object stores (*continued*)
  - moving objects to and from 245–248
  - S3 service vs. 241
  - S3 service 237
  - selecting keys for objects 256–257
  - static web hosting using
    - accessing website 254–255
    - configuring bucket 253–254
    - creating bucket for 253
  - storing objects programmatically
    - installing web application 249–250
    - listing files in bucket 251–252
    - overview 248–249
    - setting up S3 bucket 249
    - uploading files to S3 250–251
- on-demand instances 92
- OpenSwan VPN server
  - installing VPN with script 144–145
  - overview 140–142
- optimistic locking 447–448
- OS (operating system) 46, 62–64

## P

---

- PaaS (platform as a service) 5
- partition key 355–356, 358–361
- PCI (payment card industry) 12
- performance
  - database
    - increasing database resources 314–315
    - using read replication 316–318
  - increasing speed using CDN 252–253
- PermittedThroughput metric 286–287
- pessimistic locking 448
- pip tool 109
- policies, authorization 174–176
- PostgreSQL 339
- primary key, DynamoDB
  - defined 355
  - partition key 355–356
- private IP addresses 184
- public IP addresses 86–88, 184
- putItem operation 362
- PuTTY 32, 188
- Python
  - implementing Lambda function 222–223
  - installing packages 109

## Q

---

- querying data, from DynamoDB service
  - by key 365–366
  - by key and filter 366–368
  - overview 364–365

- retrieving all items 371–372
- using secondary indexes 368–371
- queue
  - decoupling dynamic EC2 instances pools
    - by 481–484
- queuing theory 345

## R

---

- RDP (Remote Desktop Protocol) 29
- RDS (relational database service)
  - access control
    - network access control 310–311
    - overview 311
    - using IAM service 309–310
  - backup/restore
    - automated snapshots 303–304
    - copying database to other region 307
    - costs 308
    - manual snapshots 304–305
    - restoring database 305–307
  - defined 37
  - DynamoDB service vs. 353
  - failures possible for 432
  - high availability 311–314
  - importing data 300–303
  - monitoring database 318–320
  - MySQL databases
    - costs 300
    - database instance information 299–300
    - WordPress platform using 297–299
  - overview 294–296
  - performance
    - increasing database resources 314–315
    - using read replication 316–318
- Redis
  - cluster with cluster mode disabled 332
  - cluster with cluster mode enabled 332–334
  - replication, whether up-to-date 345–346
  - single-node cluster 331–332
  - SortedSet 326
  - testing cluster 328–330
- regions 39, 392
- reliability, of AWS 11
- reserved instances 92–94
- resource groups 44
- REST API 103–104
- roles, authentication 177–179
- root user 172–174
- RPO (recovery point objective) 411–412
- RTO (recovery time objective) 411–412



## S

- S3 (Simple Storage Service)
  - backing up data using 238–240
  - data consistency 255
  - defined 3
  - deployment 455–457
  - Glacier service 241
  - linking custom domain to bucket 255
  - overview 237
  - selecting keys for objects 256–257
  - static web hosting using
    - accessing website 254–255
    - configuring bucket 253–254
    - creating bucket for 253
  - storing objects programmatically
    - installing web application 249–250
    - listing files in bucket 251–252
    - overview 248–249
    - setting up S3 bucket 249
    - uploading files to S3 250–251
  - versioning for 239–240
- SaaS (software as a service) 5
- SAM (Serverless Application Model)
  - defining Lambda function with 224
  - deploying Lambda function with 226
  - setting up Lambda function with 223–224
- scaling
  - advantages of AWS 11
  - based on CPU load 474
  - decoupling dynamic EC2 instance pool
    - by load balancer 476–480
    - by queue 481–484
    - overview 475–482
  - DynamoDB service 375–379
  - general discussion 463–465
  - managing dynamic EC2 instance pool 465–469
  - policies 465
  - triggering auto-scaling
    - based on CloudWatch metrics 472–474
    - based on schedule 471–472
    - overview 469–470
- scan operation 371–372
- schedule expressions 207
- scheduled events 203, 208
- SDKs (software development kits)
  - overview 21
  - platform and language support 117
- secondary indexes 368–371
- security
  - AWS account
    - authentication roles 177–179
    - authorization policies 174–176
    - creating users 176–177
    - IAM service 173–174
    - importance of securing 171–177
    - root user 172–173
  - controlling network traffic
    - allowing ICMP traffic 182–183
    - allowing SSH traffic 183–184
    - allowing SSH traffic from IP address 184–185
    - allowing SSH traffic from security group 185–189
    - overview 179–181
    - using security groups 181–182
  - creating VPC
    - accessing internet via NAT gateway 196–197
    - adding private Apache web server subnet 194–195
    - creating IGW 190–191
    - defining public bastion host subnet 192–193
    - launching virtual machines in subnets 195–196
    - overview 189–190
  - shared responsibility with AWS 167
  - updating software
    - checking for security updates 168–169
    - installing updates on running virtual machines 170–171
    - installing updates on startup 169–170
- security groups 338–342
  - allowing SSH traffic 185–189
  - defined 38
  - overview 181–182
- sharing data volumes between machines. *See* EFS (Elastic File System)
- single point of failure (SPOF) 332
- single-node cluster, Redis 331–332
- snapshots, database
  - automated 303–304
  - copying automated as manual 305
  - manual 304–305
- software, installing on virtual machines 75
- SPOF (single point of failure) 332, 432, 434–436
- spot instances 92, 95–100
- SQS (Simple Queue Service)
  - consuming messages 425–428
  - creating queue 423
  - creating worker for consuming messages 452–455
  - defined 415
  - deployment 455–457
  - limitations of 428–429
  - sending messages to queue 423–425
- SSH traffic, allowing 183–184
  - from IP address 184–185
  - from security group 185–189
- stacks 128, 152
- stateless servers 249

- statelessness 465
- static web hosting
  - accessing website 254–255
  - configuring bucket 253–254
  - creating bucket 253
- stopping vs. terminating machines 78
- storage
  - instance stores
    - backups 272
    - overview 268–271
    - performance testing 272
    - viewing and mounting volumes 271–272
  - network-attached storage
    - backing up data 266–268
    - creating volumes 260–261
    - managing volumes 261–263
    - performance improvements 263
- streams, DynamoDB 364–365
- strongly consistent reads 373, 375
- subnet group 338–340
- subnets 436
- synchronous decoupling
  - overview 415–416
  - setting up load balancer 416–420
- system status checks 385

## T

---

- tables, DynamoDB
  - creating 358
  - overview 354–355
  - using partition key 358–359
  - using partition key and sort key 360–361
- tags, defined 67
- templates, CloudFormation
  - AWSTemplateFormatVersion value 123
  - example of 127–133
  - outputs structure 126
  - overview 121–122
  - parameters structure 123–125
  - resources structure 125–126
- terminating vs. stopping machines 78
- Terraform 133
- TiB (tebibyte) 265
- tools
  - blueprints 22
  - CLI 20
  - Management Console 19
  - SDKs 21
- Troposphere 133
- TTL (time to live) 415

## U

---

- updates, security
  - checking for 168–169
  - installing on running virtual machines 170–171
  - installing on startup 169–170
- use cases
  - fault-tolerant systems 8–10
  - running Java EE applications 7–8
  - running web shop 5–7
- user data 140
- users, creating 176–177
- UUID (universally unique identifier) 438

## V

---

- versioning
  - for applications 146
  - for S3 buckets 239–240
  - versions defined 457
- virtual appliances 63–64
- virtual machines 37
  - allocating fixed public IP address 86–88
  - changing size of 79–82
  - comparing AWS Lambda with 202–203
  - connecting to
    - from Linux 73
    - from Mac OS X 73
    - from Windows 73
    - login message when connecting 74–75
    - overview 72–75
  - cost optimization
    - overview 92
    - reserved virtual machines 93–94
    - spot instances 95–100
  - creating additional network interface for 88–92
  - creating using CLI script 113–115
  - installing software on 75
  - launching
    - choosing size of 64–65
    - naming 66
    - overview 60–61
    - selecting key pair for SSH 70–72
    - selecting OS 62–64
  - listing EC2 instances using CLI 113
  - monitoring
    - load 77–78
    - viewing logs 76–77
  - overview 60–75
  - running script on virtual machine startup
    - application update process 145
    - overview 139
    - using user data 140

virtual machines (*continued*)

- security updates for
  - checking 168–169
  - installing on running virtual machines 170–171
  - installing on startup 169–170
- shutting down 78–79
- starting in another data center 82–86

## virtualization, granularity of 202

## VMs (virtual machines) 7

*See also* EC2 Instance

- VPC (Virtual Private Cloud) 337–338, 386, 395
  - accessing endpoints within 217–218
  - accessing internet via NAT gateway 196–197
  - adding private Apache web server subnet 194–195
  - creating IGW 190–191
  - defined 432
  - defining public bastion host subnet 192–193
  - launching virtual machines in subnets 195–196
  - overview 189–190
- VPN (Virtual Private Network) 7
  - installing VPN with script 144–145
  - overview 140–142
  - using CloudFormation to start virtual machine with user data 142–144

**W**

---

- web applications
  - using AWS Elastic Beanstalk
    - components of 146
    - creating application 147
- website health check
  - building 203–218
- wildcard character ( \* ) 174
- Windows
  - connecting to virtual machines from 73
  - EC instances on 270
  - installing CLI on 109–110
  - SSH client on 32
- WordPress
  - AWS installation example
    - costs 52–53
    - creating infrastructure 37–44
    - load balancer 47–48
    - MySQL database 49–51
    - web servers 45–47

**X**

---

## Xen 64

**Y**

---

- yum package manager 85–86, 144, 168

## RELATED MANNING TITLES




### *Learn Amazon Web Services in a Month of Lunches*

by David Clinton

ISBN: 9781617294440

328 pages, \$39.99

August 2017



### *AWS Lambda in Action*

*Event-driven serverless applications*

by Danilo Poccia

ISBN: 9781617293719

384 pages, \$49.99

November 2016




### *Google Cloud Platform in Action*

by JJ Geewax

ISBN: 9781617293528

632 pages, \$59.99

September 2018



### *Learn Azure in a Month of Lunches*

by Iain Foulds

ISBN: 9781617295171

375 pages, \$44.99

September 2018

*For ordering information go to [www.manning.com](http://www.manning.com)*

## Storage & Database

Abbr.	Name	Description	Where
S3	Amazon Simple Storage Service	Object store to save data without any size restrictions	8
	Amazon Glacier	Inexpensive data archive solution	8.4
EBS	Amazon Elastic Block Store	Network attached block level storage for EC2 instances	9.1
	Amazon EC2 Instance Store	Block level storage for EC2 instances	9.2
EFS	Amazon Elastic File System	Scalable network file system based on NFSv4	10
RDS	Amazon Relational Database Service	MySQL, Oracle Database, Microsoft SQL Server, or PostgreSQL	11
	Amazon DynamoDB	Proprietary NoSQL key-value database with limited but powerful options for queries	13
	Amazon ElastiCache	In-memory data store based on Redis or Memcached typically used to cache data	12

## Architecting on AWS

Abbr.	Name	Description	Where
AZ	Availability Zone	Group of isolated data centers within a region	14.2
ASG	Amazon EC2 auto-scaling group	Observes a fleet of servers: Replaces faulty servers and increases/decreases the size of the fleet based on external triggers like CPU usage	17
	Amazon CloudWatch	Keeps track of metrics and triggers alarms when a threshold is reached	14.1
ELB	Elastic Load Balancing	Load balancer for EC2 instances	15.1
ALB	Application Load Balancer	Layer 7 load balancer with HTTP and HTTPS support	15.1
SQS	Amazon Simple Queue Service	Message queue	15.2

# Amazon Web Services IN ACTION Second Edition

Michael Wittig • Andreas Wittig

**T**he largest and most mature of the cloud platforms, AWS offers over 100 prebuilt services, practically limitless compute resources, bottomless secure storage, as well as top-notch automation capabilities. This book shows you how to develop, host, and manage applications on AWS.

**Amazon Web Services in Action, Second Edition** is a comprehensive introduction to deploying web applications in the AWS cloud. You'll find clear, relevant coverage of all essential AWS services, with a focus on automation, security, high availability, and scalability. This thoroughly revised edition covers the latest additions to AWS, including serverless infrastructure with AWS Lambda, sharing data with EFS, and in-memory storage with ElastiCache.

## What's Inside

- Completely revised bestseller!
- Secure and scale distributed applications
- Deploy applications on AWS
- Design for failure to achieve high availability
- Automate your infrastructure

Written for mid-level developers and DevOps engineers.

**Andreas** and **Michael Wittig** are software engineers and DevOps consultants focused on AWS. Together, they migrated the first bank in Germany to AWS in 2013.

To download their free eBook in PDF, ePub, and Kindle formats, owners of this book should visit  
[manning.com/books/amazon-web-services-in-action-second-edition](http://manning.com/books/amazon-web-services-in-action-second-edition)

“Slices through the complexity of AWS using examples and visuals to cement knowledge in the minds of readers.”

—From the Foreword by  
Ben Whaley  
AWS community hero and author

“The authors’ ability to explain complex concepts is the real strength of the book.”

—Antonio Pessolano  
Consoft Sistemi

“Useful examples, figures, and sources to help you learn efficiently.”

—Christof Marte, Daimler-Benz

“Does a great job of explaining some of the key services in plain English so you have the knowledge necessary to dig deeper.”

—Ryan Burrows  
Rooster Park Consulting



ISBN-13: 978-1-61729-511-9  
 ISBN-10: 1-61729-511-6



9 781617 295119



MANNING

\$54.99 / Can \$72.99 [INCLUDING eBook]